

CSP 554 Big Data Technologies

Assignment #4

Student ID: A20549927

Prachi Kotadia

To execute the TestDataGen file to obtain the magic number 11895


```
[hadoop@ip-172-31-58-223 ~]$ ls
TestDataGen.class
[hadoop@ip-172-31-58-223 ~]$ java TestDataGen
Magic Number = 11895
```

To copy foodratings11895.txt and foodplaces11895.txt file into /user/hadoop:

```
[hadoop@ip-172-31-58-223 ~]$ hadoop fs -copyFromLocal /home/hadoop/foodratings11895.txt /user/hadoop
[hadoop@ip-172-31-58-223 ~]$ hadoop fs -copyFromLocal /home/hadoop/foodplaces11895.txt /user/hadoop
[hadoop@ip-172-31-58-223 ~]$ hadoop fs -ls /user/hadoop
Found 2 items
-rw-r--r-- 1 hadoop hdfsadmingroup 59 2024-03-07 16:16 /user/hadoop/foodplaces11895.txt
-rw-r--r-- 1 hadoop hdfsadmingroup 17476 2024-03-07 16:16 /user/hadoop/foodratings11895.txt
```

For all exercises, load pyspark:

```
[hadoop@ip-172-31-58-223 ~]$ pyspark
Python 3.9.16 (main, Sep 8 2023, 00:00:00)
[GCC 11.4.1 20230605 (Red Hat 11.4.1-2)] on linux
Type "help", "copyright", "credits" or "license" for more information.
Mar 07, 2024 4:48:22 PM org.apache.spark.launcher.Log4jHotPatchOption staticJavaAgentOptio
WARNING: spark.log4jHotPatch.enabled is set to true, but /usr/share/log4j-cve-2021-44228-h
^[[Asetting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel)
24/03/07 16:48:25 WARN HiveConf: HiveConf of name hive.server2.thrift.url does not exist
24/03/07 16:48:27 WARN Client: Neither spark.yarn.jars nor spark.yarn.archive is set, fall
welcome to

 version 3.5.0-amzn-0

Using Python version 3.9.16 (main, Sep 8 2023 00:00:00)
Spark context web UI available at http://ip-172-31-58-223.ec2.internal:4040
Spark context available as 'sc' (master = yarn, app id = application_1709825770370_0002).
SparkSession available as 'spark'.
```

Exercise 1:

To create the foodratings dataframe:

```
from pyspark.sql.types import StructType, StructField, StringType, IntegerType

foodratings_schema = StructType([StructField("name", StringType(), True), StructField("food1",
IntegerType(), True), StructField("food2", IntegerType(), True), StructField("food3", IntegerType(), True),
StructField("food4", IntegerType(), True), StructField("placeid", IntegerType(), True)])
```

To load the foodratings11895.txt into foodratings dataframe:

```
foodratings =
spark.read.format("csv").schema(foodratings_schema).load("hdfs:///user/hadoop/foodratings11895.txt")
```

To load the foodratings11895.txt into foodratings dataframe: foodratings =

```
spark.read.format("csv").schema(foodratings_schema).load("hdfs:///user/hadoop/foodratings11895.txt")
```

To print the schema from foodratings dataframe:

```
foodratings.printSchema()
```

To print the top 5 rows from foodratings dataframe:

```
foodratings.show(5)
```

```
>>> from pyspark.sql.types import StructType, StructField, StringType, IntegerType
>>> foodratings_schema = StructType([StructField("name", StringType(), True), StructField("food1", IntegerType(), True), StructField("food2", IntegerType(), True), StructField("food3", IntegerType(), True), StructField("food4", IntegerType(), True), StructField("placeid", IntegerType(), True)])
>>> foodratings = spark.read.format("csv").schema(foodratings_schema).load("hdfs:///user/hadoop/foodratings11895.txt")
>>> foodratings.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)

>>> foodratings.show(5)
+-----+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+-----+-----+-----+-----+-----+
|Mel|13|26|26|18|5|
|Sam|28|35|37|39|1|
|Sam|21|50|25|36|1|
|Mel|47|7|16|23|2|
|Joe|17|25|41|45|1|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Exercise 2:

To create foodplaces dataframe: `foodplaces_schema = StructType([StructField("placeid", IntegerType(), True), StructField("placename", StringType(), True)])`

```
>>> foodplaces_schema = StructType([StructField("placeid", IntegerType(), True), StructField("placename",
... StringType(), True)])
```

To load foodplaces11895.txt into foodplaces dataframe:

```
foodplaces =
park.read.format("csv").schema(foodplaces_schema).load("hdfs:///user/hadoop/foodplaces11895.txt")
```

To print the foodplaces dataframe schema:

```
foodplaces.printSchema()
```

To print the top 5 rows from foodplaces dataframe:

```
foodplaces.show(5)
```

```
>>> foodplaces = spark.read.format("csv").schema(foodplaces_schema).load("hdfs:///user/hadoop/foodplaces11895.txt")
>>> foodplaces.printSchema()
root
 |-- placeid: integer (nullable = true)
 |-- placename: string (nullable = true)

>>> foodplaces.show(5)
+-----+-----+
|placeid|placename|
+-----+-----+
|1|China Bistro|
|2|Atlantic|
|3|Food Town|
|4|Jake's|
|5|Soup Bowl|
+-----+-----+
```

Exercise 3:

Step A:

To register the dataframes from Exercise 1 and Exercise called “foodratingsT” and “foodplacesT”:

```
foodratings.createOrReplaceTempView("foodratingsT")
```

```
foodplaces.createOrReplaceTempView("foodplacesT")
```

```
>>> foodratings.createOrReplaceTempView("foodratingsT")
>>> foodplaces.createOrReplaceTempView("foodplacesT")
```

Step B:

To create foodratings_ex3a dataframe to print the schema and print the top 5 rows from that dataframe:

```
foodratings_ex3a = spark.sql("select * from foodratingsT where food2 < 25 and food4 > 40")
```

```
foodratings_ex3a.printSchema() foodratings_ex3a.show(5)
```

```
>>> foodratings_ex3a = spark.sql("select * from foodratingsT where food2 < 25 and food4 > 40")
>>> foodratings_ex3a.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)

>>> foodratings_ex3a.show(5)
+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+-----+-----+-----+-----+-----+
| Joy|   46|   15|   47|   49|      3|
| Joy|   30|   21|   32|   42|      5|
| Joe|   47|    6|   43|   50|      2|
| Joe|   27|   17|   44|   50|      5|
| Joy|   14|   20|   50|   49|      3|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Step C:

To create foodratings_ex3b dataframe to print the schema and print the top 5 rows from that dataframe:

```
foodplaces_ex3b = spark.sql("select * from foodplacesT where placeid > 3")
```

```
foodplaces_ex3b.printSchema() foodplaces_ex3b.show(5)
```

```
>>> foodplaces_ex3b = spark.sql("select * from foodplacesT where placeid > 3")
>>> foodplaces_ex3b.printSchema()
root
 |-- placeid: integer (nullable = true)
 |-- placename: string (nullable = true)

>>> foodplaces_ex3b.show(5)
+-----+-----+
|placeid|placename|
+-----+-----+
|      4|   Jake's|
|      5| Soup Bowl|
+-----+-----+
```

Exercise 4:

To create `foodratings_ex4` dataframe using transformation to print the schema and print the top 5 rows from that dataframe:

```
foodratings_ex4 = foodratings.filter((foodratings['name'] == "Mel") & (foodratings['food3'] < 25))
```

```
foodratings_ex4.printSchema() foodratings_ex4.show(5)
```

```
>>> foodratings_ex4 = foodratings.filter((foodratings['name'] == "Mel") & (foodratings['food3'] < 25))
>>> foodratings_ex4.printSchema()
root
|-- name: string (nullable = true)
|-- food1: integer (nullable = true)
|-- food2: integer (nullable = true)
|-- food3: integer (nullable = true)
|-- food4: integer (nullable = true)
|-- placeid: integer (nullable = true)

>>> foodratings_ex4.show(5)
+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+-----+-----+-----+-----+-----+
| Mel|   47|    7|   16|   23|     2|
| Mel|    9|   21|   17|   33|     2|
| Mel|   23|   42|   13|   37|     4|
| Mel|   25|   50|   24|   44|     1|
| Mel|   10|   46|    6|    2|     2|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Exercise 5:

To create `foodratings_ex5` dataframe using transformation to print the schema and print the top 5 rows from that dataframe:

```
foodratings_ex5 = foodratings.select('name', 'placeid')
```

```
foodratings_ex5.printSchema() foodratings_ex5.show(5)
```

```
>>> foodratings_ex5 = foodratings.select('name', 'placeid')
>>> foodratings_ex5.printSchema()
root
|-- name: string (nullable = true)
|-- placeid: integer (nullable = true)

>>> foodratings_ex5.show(5)
+-----+-----+
|name|placeid|
+-----+-----+
| Mel|      5|
| Sam|      1|
| Sam|      1|
| Mel|      2|
| Joe|      1|
+-----+-----+
only showing top 5 rows
```

Exercise 6:

To create ex6 dataframe using transformation to print the schema and print the top 5 rows from that dataframe:

```
ex6 = foodratings.join(foodplaces, foodratings.placeid == foodplaces.placeid, 'inner')
ex6.printSchema() ex6.show(5)
```

```
>>> ex6 = foodratings.join(foodplaces, foodratings.placeid == foodplaces.placeid, 'inner')
>>> ex6.printSchema()
root
|-- name: string (nullable = true)
|-- food1: integer (nullable = true)
|-- food2: integer (nullable = true)
|-- food3: integer (nullable = true)
|-- food4: integer (nullable = true)
|-- placeid: integer (nullable = true)
|-- placeid: integer (nullable = true)
|-- placename: string (nullable = true)

>>> ex6.show(5)
+---+---+---+---+---+---+---+
|name|food1|food2|food3|food4|placeid|placeid|  placename|
+---+---+---+---+---+---+---+
| Mel |   13 |   26 |   26 |   18 |     5 |     5 |  Soup Bowl |
| Sam |   28 |   35 |   37 |   39 |     1 |     1 | china Bistro |
| Sam |   21 |   50 |   25 |   36 |     1 |     1 | china Bistro |
| Mel |   47 |    7 |   16 |   23 |     2 |     2 |   Atlantic |
| Joe |   17 |   25 |   41 |   45 |     1 |     1 | china Bistro |
+---+---+---+---+---+---+---+
only showing top 5 rows
```

Terminate the Cluster

My cluster_assign7

Updated less than a minute ago

Terminate

Clone in AWS CLI

Clone

▼ Summary

Cluster info

Cluster ID
j-3W0V15UCTN6E

Cluster configuration
Instance groups

Capacity
1 Primary 1 Core 0 Task

Applications

Amazon EMR version
emr-7.0.0

Installed applications
Hadoop 3.3.6, Hive 3.1.3,
JupyterEnterpriseGateway 2.6.0, Livy 0.7.1,
Spark 3.5.0

Cluster management

Log destination in Amazon S3
[aws-logs-905418339338-us-east-1/elasticmapreduce](#)

Persistent application UIs
[Spark History Server](#)
[YARN timeline server](#)
[Tez UI](#)

Status and time

Status
[Terminating](#)

Creation time
March 07, 2024, 09:30 (UTC-06:00)

Elapsed time
1 hour, 59 minutes